Documentation 1B:

1. Introduction:

In this analysis, we explore the application of Principal Component Analysis (PCA) and linear regression to predict salary values in a baseball dataset. The code follows a structured approach, including data loading, preprocessing, exploratory data analysis, model implementation, and results analysis.

Here is the following data description:

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	\
0	293	66	1	30	29	14	1	293	66	1	30	
1	315	81	7	24	38	39	14	3449	835	69	321	
2	479	130	18	66	72	76	3	1624	457	63	224	
3	496	141	20	65	78	37	11	5628	1575	225	828	
4	321	87	10	39	42	30	2	396	101	12	48	

	CRBI	CWalks	PutOuts	Assists	Errors	Salary
0	29	14	446	33	20	NaN
1	414	375	632	43	10	475.0
2	266	263	880	82	14	480.0
3	838	354	200	11	3	500.0
4	46	33	805	40	4	91.5

	AtBat	Hits	HmRun	Runs	RBI	Walks	\
count	322.000000	322.000000	322.000000	322.000000	322.000000	322.000000	
mean	380.928571	101.024845	10.770186	50.909938	48.027950	38.742236	
std	153.404981	46.454741	8.709037	26.024095	26.166895	21.639327	
min	16.000000	1.000000	0.000000	0.000000	0.000000	0.000000	
25%	255.250000	64.000000	4.000000	30.250000	28.000000	22.000000	
50%	379.500000	96.000000	8.000000	48.000000	44.000000	35.000000	
75%	512.000000	137.000000	16.000000	69.000000	64.750000	53.000000	
max	687.000000	238.000000	40.000000	130.000000	121.000000	105.000000	

2. Data Loading and Preprocessing

The dataset is loaded from an Excel file (Hitters.xlsx), and categorical columns are dropped. To handle missing values, entries with NaN in the 'Salary' feature are removed.

Here in the below we can see that Salary has NAN:

	CRBI	CWalks	PutOuts	Assists	Errors	Salary
0	29	14	446	33	20	NaN
1	414	375	632	43	10	475.0
2	266	263	880	82	14	480.0
3	838	354	200	11	3	500.0
4	46	33	805	40	4	91.5

The number of NAN are depicted below :

```
AtBat       0
Hits        0
HmRun       0
Runs        0
RBI         0
Walks       0
Years       0
CAtBat      0
CHits       0
CHmRun      0
CRuns       0
CRBI        0
CWalks      0
PutOuts     0
Assists     0
Errors      0
Salary     59
dtype: int64
```

3. Exploratory Data Analysis (EDA)

Basic statistics and a pair plot are generated to gain insights into the dataset's structure and relationships between variables.

Here is the sns pair plot :

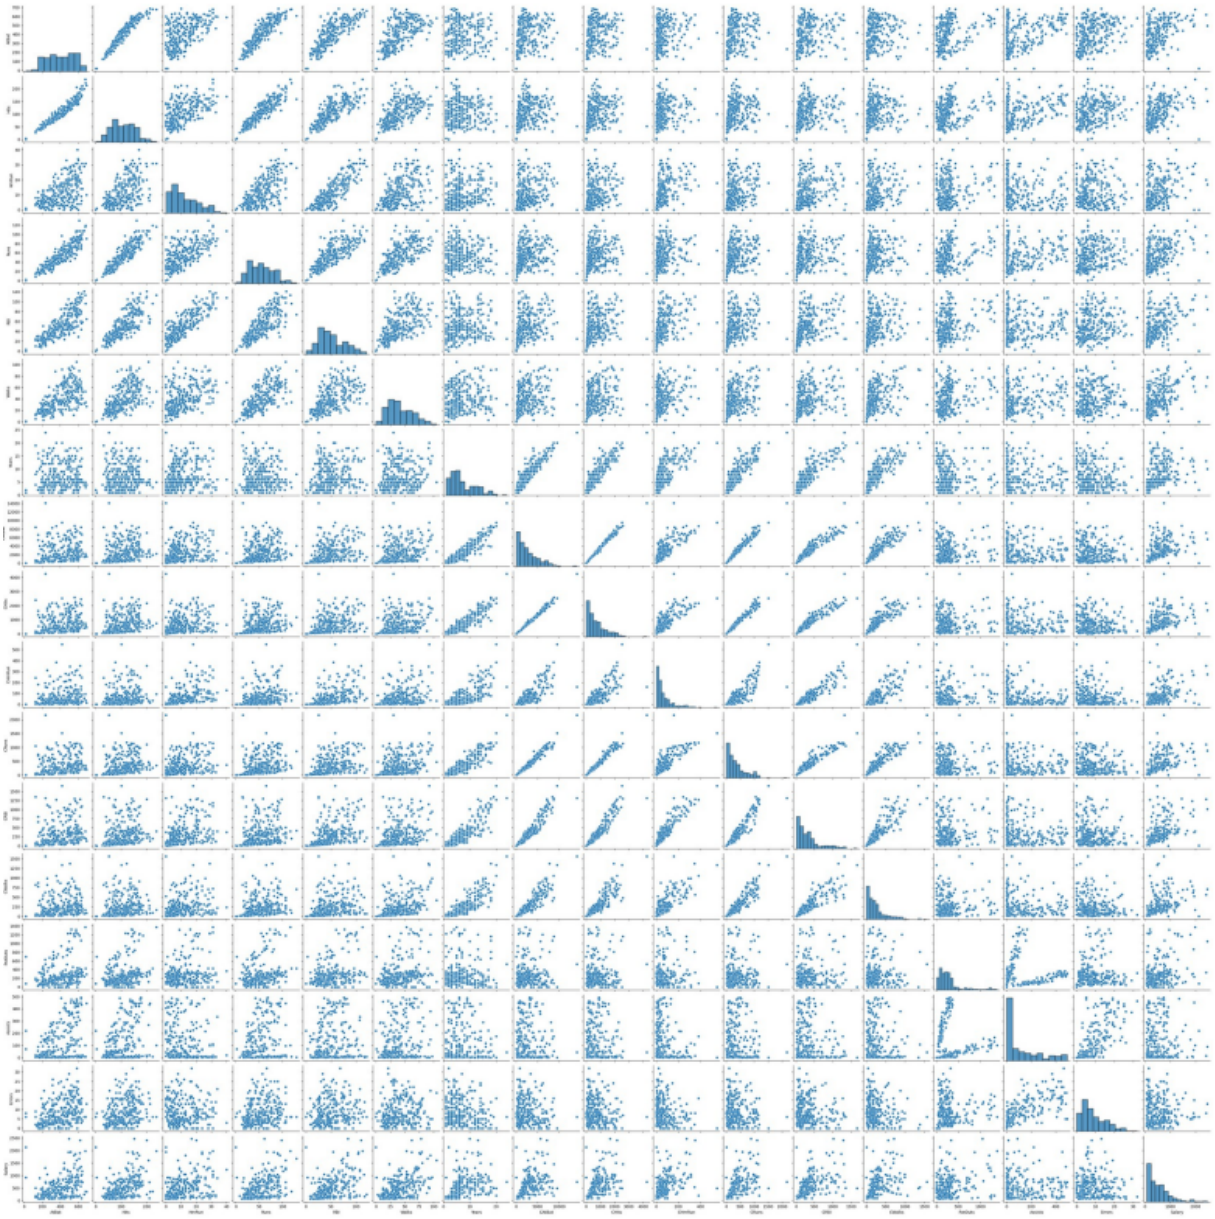

4. Principal Component Analysis (PCA)

- Features are standardized to ensure uniformity.

- The covariance matrix, eigenvalues, and eigenvectors are calculated.
- Eigenvalues and eigenvectors are sorted to determine the principal components.
- Explained variance ratio and cumulative explained variance are computed.
- The number of components necessary for efficient prediction is identified through cumulative explained variance.

Here is the plot to determine the establishment of number of components vs Cumulative Variance

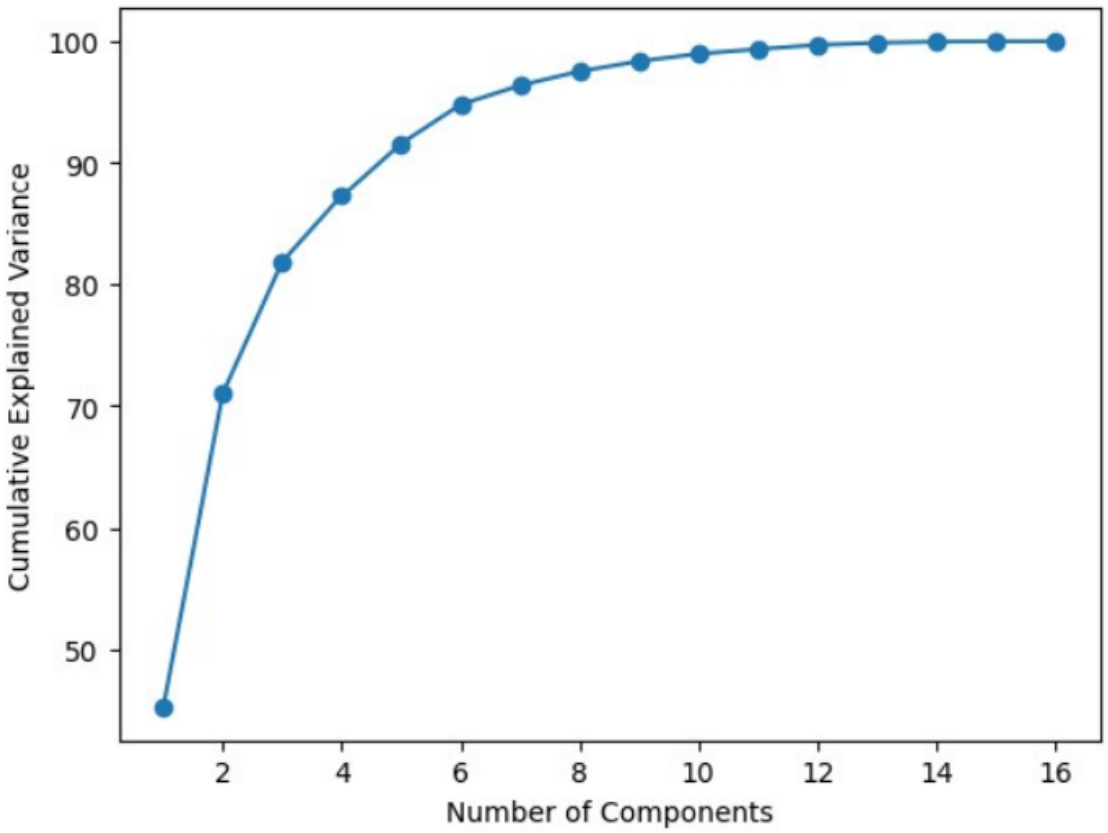

5. Number of Components vs RMSE

A loop iterates over different numbers of principal components, and for each iteration:

- Training and test data are transformed using PCA.
- Linear regression is applied.

- Root Mean Squared Error (RMSE) is calculated for each iteration.

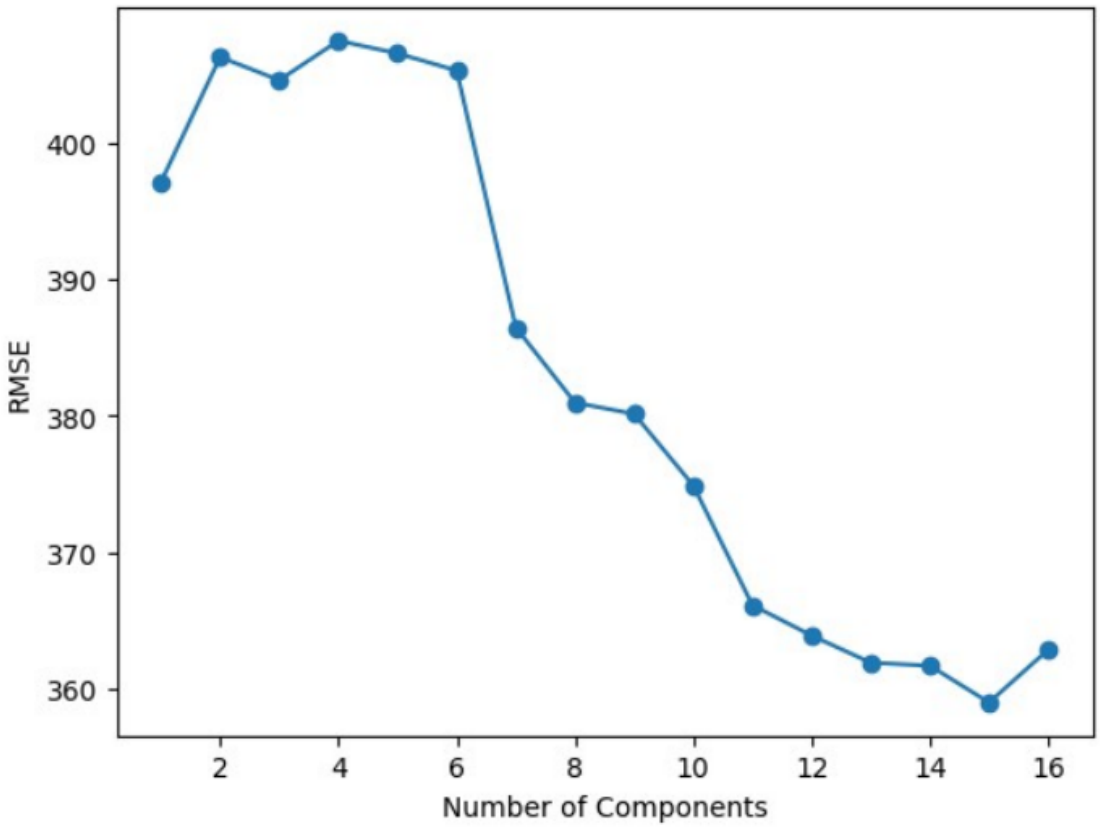

6. Results Analysis

- The cumulative explained variance and RMSE vs number of components are presented graphically.

  The optimal number of components is selected based on minimal RMSE.

	Num of Principal Components	RMSE
0	1	397.135817
1	2	406.284985
2	3	404.594695
3	4	407.488046
4	5	406.552058
5	6	405.300602
6	7	386.470377
7	8	380.968342
8	9	380.142136
9	10	374.898270
10	11	366.128983
11	12	363.906810
12	13	361.919565
13	14	361.701757
14	15	359.003059
15	16	362.870206

- From the above table we have decided to take the optimal number of components but there might be a question of not taking into account lower RMSE when the number of features are 12 ,13,14.

- But the reason for taking the optimal no of components to be taken as 15 is described in the below table:

	Num of Principal Components	Percentage Variance covered
0	1	45.311913
1	2	70.998464
2	3	81.798318
3	4	87.238195
4	5	91.596762
5	6	94.799095
6	7	96.372127
7	8	97.528009
8	9	98.355667
9	10	98.967987
10	11	99.351456
11	12	99.695435
12	13	99.871587
13	14	99.961897
14	15	99.992487
15	16	100.000000

- Since the encapsulated variance is maximum in case of 15 features and RMSE being close enough wrt to the number of features being 12 ,we therefore consider 15 as the optimal number of features.

- Upon mounting Linear Regression model and running it through the testing data here is the tabulation of Actual vs Predicted Values:

	Actual Values	Predicted Values
154	277.5	344.055982
279	150.0	951.286871
221	210.0	128.226948
7	100.0	987.963759
307	277.5	-250.433812
17	175.0	344.729344
176	86.5	85.873782

- Here is the Corresponding plot to depict the Actual vs Predicted values :